

Machine Learning - Logistic Regression, Linear Discriminant and Quadratic Discriminant Analysis (Titanic data set)

```
In [1]: #Import of packages  
import numpy as np
```

```
In [2]: import pandas as pd
```

```
In [3]: import matplotlib.pyplot as plt #for graphing
```

```
In [4]: import seaborn as sb #for graphing
```

```
In [5]: # Import of Scikit-Learn  
import sklearn
```

```
In [6]: from sklearn import preprocessing
```

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [7]: from sklearn.linear_model import LogisticRegression
```

```
In [9]: from sklearn import metrics
```

```
In [10]: from sklearn.metrics import classification_report
```

```
In [11]: from sklearn.metrics import accuracy_score
```

```
In [12]: %matplotlib inline
```

```
In [ ]:
```

Read-in data, analyse data

```
In [13]: train = pd.read_csv('C:/Bilder/train.csv')
```

```
In [14]: train.head()
```

```
Out[14]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Ca
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	1
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	1
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	1
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	1

```
In [15]: # Rows and columns
train.shape
```

```
Out[15]: (891, 12)
```

```
In [16]: # Descriptive statistics
train.describe()
```

```
Out[16]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [17]: train.dtypes
```

```
Out[17]: PassengerId    int64
Survived      int64
Pclass        int64
Name          object
Sex           object
Age           float64
SibSp         int64
Parch         int64
Ticket        object
Fare          float64
Cabin         object
Embarked      object
dtype: object
```

```
In [18]: # Find out how many missing values are there
train.isnull().sum()
```

```
Out[18]: PassengerId    0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                 177
SibSp                0
Parch               0
Ticket              0
Fare                0
Cabin               687
Embarked            2
dtype: int64
```

Pre-processing

Survived is y-variable

```
In [19]: train.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
train.head()
```

```
Out[19]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

```
In [20]: train.Pclass.value_counts()
```

```
Out[20]: Pclass
3      491
1      216
2      184
Name: count, dtype: int64
```

```
In [21]: train.Pclass = train.Pclass.apply(str)
```

```
In [22]: train.dtypes
```

```
Out[22]: Survived      int64
Pclass      object
Sex         object
Age        float64
SibSp      int64
Parch      int64
Fare       float64
Embarked   object
dtype: object
```

```
In [23]: train = pd.get_dummies(train, prefix_sep='_', drop_first=True, dtype=float)
```

```
In [24]: train.head()
```

```
Out[24]:
```

	Survived	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male	Embarked_Q	Embarke
0	0	22.0	1	0	7.2500	0.0	1.0	1.0	0.0	
1	1	38.0	1	0	71.2833	0.0	0.0	0.0	0.0	
2	1	26.0	0	0	7.9250	0.0	1.0	0.0	0.0	
3	1	35.0	1	0	53.1000	0.0	0.0	0.0	0.0	
4	0	35.0	0	0	8.0500	0.0	1.0	1.0	0.0	

```
In [25]: train.isnull().sum()
```

```
Out[25]: Survived      0
Age          177
SibSp        0
Parch        0
Fare         0
Pclass_2     0
Pclass_3     0
Sex_male     0
Embarked_Q   0
Embarked_S   0
dtype: int64
```

In [26]: train.dtypes

```
Out[26]: Survived      int64
Age          float64
SibSp        int64
Parch        int64
Fare         float64
Pclass_2     float64
Pclass_3     float64
Sex_male     float64
Embarked_Q   float64
Embarked_S   float64
dtype: object
```


In [27]: train.corr()

```
Out[27]:
```

	Survived	Age	SibSp	Parch	Fare	Pclass_2	Pclass_3	Sex_male
Survived	1.000000	-0.077221	-0.035322	0.081629	0.257307	0.093349	-0.322308	-0.543351
Age	-0.077221	1.000000	-0.308247	-0.189119	0.096067	0.006954	-0.312271	0.093254
SibSp	-0.035322	-0.308247	1.000000	0.414838	0.159651	-0.055932	0.092548	-0.114631
Parch	0.081629	-0.189119	0.414838	1.000000	0.216225	-0.000734	0.015790	-0.245489
Fare	0.257307	0.096067	0.159651	0.216225	1.000000	-0.118557	-0.413333	-0.182333
Pclass_2	0.093349	0.006954	-0.055932	-0.000734	-0.118557	1.000000	-0.565210	-0.064746
Pclass_3	-0.322308	-0.312271	0.092548	0.015790	-0.413333	-0.565210	1.000000	0.137143
Sex_male	-0.543351	0.093254	-0.114631	-0.245489	-0.182333	-0.064746	0.137143	1.000000
Embarked_Q	0.003650	-0.022405	-0.026354	-0.081228	-0.117216	-0.127301	0.237449	-0.074111
Embarked_S	-0.155660	-0.032523	0.070941	0.063036	-0.166603	0.192061	-0.009511	0.125722

```
In [28]: sb.heatmap(train.corr(), cmap='Blues')
```

```
Out[28]: <Axes: >
```



Logistic Regression

```
In [ ]:
```

```
In [29]:
```

```
X = train  
X = train.dropna()
```

```
In [30]: X.isnull().sum()
```

```
Out[30]: Survived      0  
Age                0  
SibSp              0  
Parch              0  
Fare               0  
Pclass_2           0  
Pclass_3           0  
Sex_male           0  
Embarked_Q         0  
Embarked_S         0  
dtype: int64
```

```
In [31]: Y = X['Survived']
```

```
In [32]: Y.head()
```

```
Out[32]: 0      0  
1      1  
2      1  
3      1  
4      0  
Name: Survived, dtype: int64
```

```
In [33]: X.pop('Survived')  
X.columns
```

```
Out[33]: Index(['Age', 'SibSp', 'Parch', 'Fare', 'Pclass_2', 'Pclass_3', 'Sex_male',  
               'Embarked_Q', 'Embarked_S'],  
              dtype='object')
```

```
In [34]: X.shape
```

```
Out[34]: (714, 9)
```

```
In [35]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = .25, ran
```

```
In [36]: Y_train.shape
```

```
Out[36]: (535,)
```

```
In [37]: Y_test.shape
```

```
Out[37]: (179,)
```

```
In [38]: X_train.shape
```

```
Out[38]: (535, 9)
```

```
In [39]: X_test.shape
```

```
Out[39]: (179, 9)
```

```
In [77]: # Model of Logistic Regression
```

```
logit = LogisticRegression()
```

```
#logit = LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercep  
#intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
#penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
#verbose=0, warm_start=False)
```



```
In [78]: # Fit the model to the training data  
logit.fit(X_train, Y_train)
```

Out[78]:

▸	LogisticRegression
---	--------------------

Predictions

```
In [43]: Y_pred = logit.predict(X_test)
```

```
In [44]: Y_pred
```

```
Out[44]: array([1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1,
                0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0,
                0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0,
                1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
                0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0,
                0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
                0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 1], dtype=int64)
```

```
In [45]: Y_pred.shape
```

```
Out[45]: (179,)
```

```
In [ ]:
```

```
In [46]: Y_test.head(20) # true y's from test data sample
```

```
Out[46]: 501    0
          315    1
          679    1
          866    1
          119    0
          254    0
          885    0
          155    0
          169    0
          683    0
          515    0
          192    1
          724    1
          393    1
          523    1
          467    0
          659    0
          663    0
          777    1
          805    0
          Name: Survived, dtype: int64
```

For model evaluation: compare Y_test with Y_pred
!

```
logit.predict_proba(X_test)
```

Model Metrics

```
Y_test.value_counts()
```

```
Survived
0      103
1       76
Name: count, dtype: int64
```

Survived 0 103 1 76 Name: count, dtype: int64

```
Y_pred.shape
```

(179,)

```
confusion_matrix = metrics.confusion_matrix(Y_test, Y_pred) # Confusion Matrix
print(confusion_matrix)
```

```
[[87 16]
 [21 55]]
```

```
In [51]: print(metrics.classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.84	0.82	103
1	0.77	0.72	0.75	76
accuracy			0.79	179
macro avg	0.79	0.78	0.79	179
weighted avg	0.79	0.79	0.79	179

```
In [52]: print(metrics.accuracy_score(Y_test, Y_pred))
```

```
0.7932960893854749
```

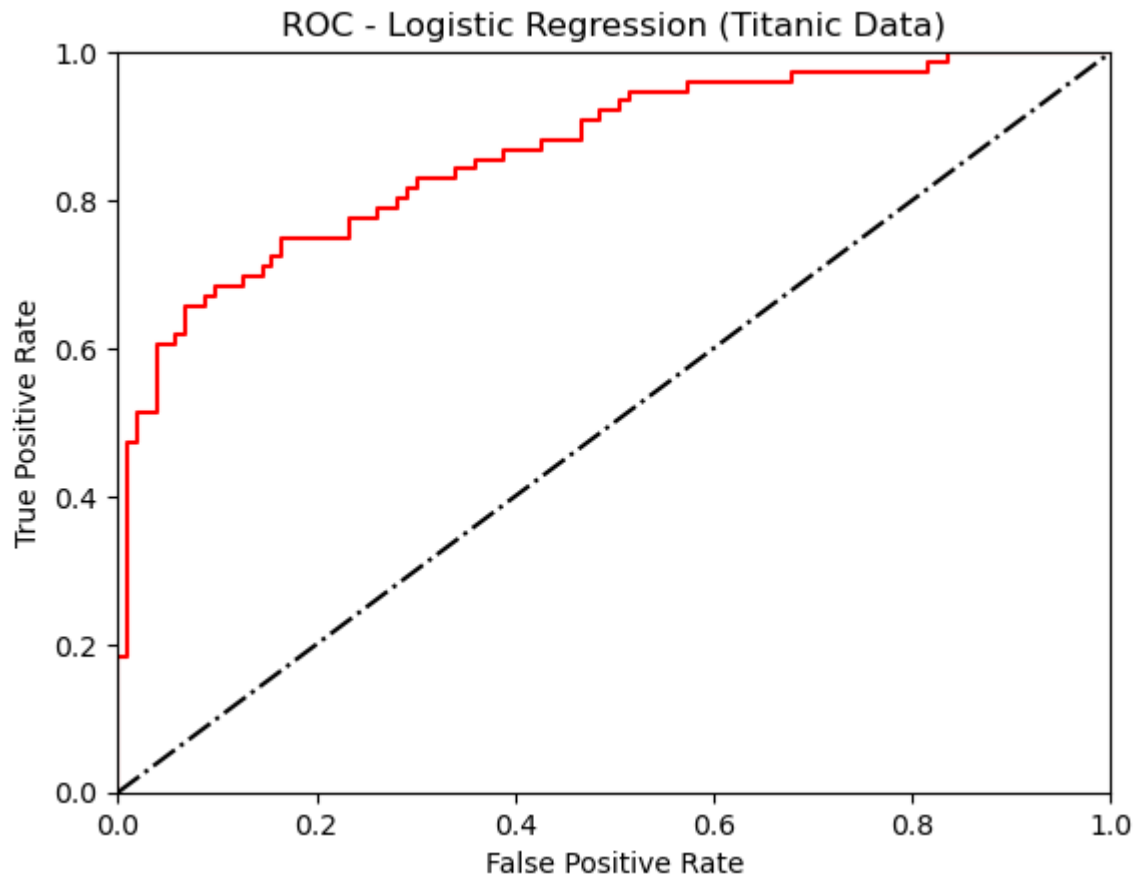
Plotting the ROC:

```
In [53]: prob = logit.predict_proba(X_test)
```

```
In [54]: pred = prob[:,1]
```

```
In [55]: fpr, tpr, threshold = metrics.roc_curve(Y_test, pred)
```

```
In [57]: plt.title('ROC - Logistic Regression (Titanic Data)')
plt.plot(fpr, tpr, 'r')
plt.plot([0, 1], [0, 1], 'k-')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [56]: metrics.roc_auc_score(Y_test, pred)
```

```
Out[56]: 0.8657383750638733
```

```
In [ ]:
```

Linear Discriminant Analysis

```
In [58]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
In [59]: lda = LinearDiscriminantAnalysis()
```

```
In [60]: lda.fit(X_train, Y_train)
```

```
Out[60]: ▼ LinearDiscriminantAnalysis
LinearDiscriminantAnalysis()
```

```
In [61]: Y_pred = lda.predict(X_test)
```

```
In [62]: Y_pred
```

```
Out[62]: array([1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1,
                0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0,
                0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0,
                1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
                0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0,
                0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
                0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 1], dtype=int64)
```

```
In [63]: Y_pred.shape
```

```
Out[63]: (179,)
```

```
In [64]: confusion_matrix = metrics.confusion_matrix(Y_test, Y_pred) # Confusion Matrix
print(confusion_matrix)
```

```
[[82 21]
 [20 56]]
```

```
In [65]: print(metrics.classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.80	0.80	103
1	0.73	0.74	0.73	76
accuracy			0.77	179
macro avg	0.77	0.77	0.77	179
weighted avg	0.77	0.77	0.77	179

```
In [66]: print(metrics.accuracy_score(Y_test, Y_pred))
```

```
0.770949720670391
```

```
In [ ]:
```

Quadratic Discriminant Analysis

```
In [67]: from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
```

```
In [68]: qda = QuadraticDiscriminantAnalysis()
```

```
In [69]: qda.fit(X_train, Y_train)
```

```
Out[69]: ▾ QuadraticDiscriminantAnalysis
          QuadraticDiscriminantAnalysis()
```

```
In [70]: Y_pred = qda.predict(X_test)
```

```
In [71]: Y_pred
```

```
Out[71]: array([1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1,
                0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0,
                0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0,
                1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
                0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0,
                0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
                0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0,
                0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 1], dtype=int64)
```

```
In [72]: Y_pred.shape
```

```
Out[72]: (179,)
```

```
In [73]: confusion_matrix = metrics.confusion_matrix(Y_test, Y_pred) # Confusion Matrix
print(confusion_matrix)
```

```
[[83 20]
 [24 52]]
```

```
In [74]: print(metrics.classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.81	0.79	103
1	0.72	0.68	0.70	76
accuracy			0.75	179
macro avg	0.75	0.75	0.75	179
weighted avg	0.75	0.75	0.75	179

```
In [75]: print(metrics.accuracy_score(Y_test, Y_pred))
```

```
0.7541899441340782
```

```
In [ ]:
```


In []:

In []: